
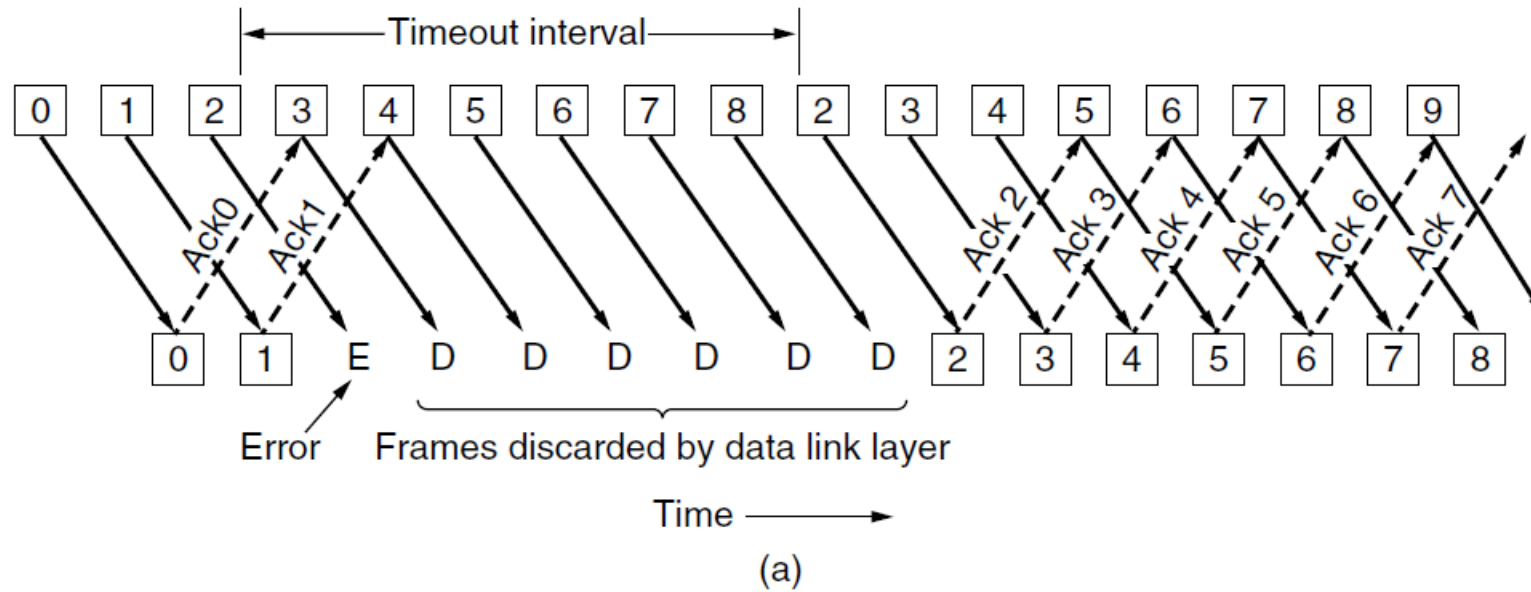


Go Back N vs. Selective Repeat



Pipelining and error recovery. Effect of an error when
(a) Go back N: receiver's window size is 1 and (b) Selective Repeat: receiver's window size is large.

Go Back N vs. Selective Repeat

- **Go Back N**

- Receiver discards all frames after lost or damaged frame
- Receiver acknowledges received frame
- Receiver does NOT send any ACK for frames received after lost or damaged frame
- Relies on sender timeout
- Equivalent to receiver window of size **1**
- Wastes a lot of bandwidth in case of error
- Receiver needs to *buffer 1 frame* only

- **Selective Repeat**

- Receiver buffers all frames received within window
- Receiver acknowledges *last received in sequence frame* for every out of sequence frame
- On timeout, sender only re-transmits *oldest unacknowledged* frame
- Receiver can deliver all buffered frames, *in sequence*, to the network layer
- Receiver often use **NAK** to stimulate retransmission before time
- Receiver needs to *buffer multiple frames* up to window size

→ *Tradeoff between buffer space and link bandwidth utilization*

Go Back N

- Allow multiple outstanding frames at Sender
 - Outstanding frame: Frame sent but not yet acknowledged
- Sender cannot send more than MAX_SEQ to
 - enforce flow control
 - avoid too much waste in case of packet loss/damage and large timeout
- *Dropped* the assumption that network layer has *infinite* supply of packets
 - Network layer causes event *network_layer_ready* when it wants to send
- Need to enforce flow control of no more than MAX_SEQ outstanding buffers at sender
 - *enable_network_layer()* to allow network layer to send
 - *disable_network_layer()* to block network layer
- On timer expire, *all outstanding frames* are re-sent

Go Back N

- Circular sequence number:

```
static boolean between(seq_nr a, seq_nr b, seq_nr c)
{
    /* Return true if a <= b < c circularly; false otherwise. */
    if (((a <= b) && (b < c)) || ((c < a) && (a <= b)) || ((b < c) && (c < a)))
        return(true);
    else
        return(false);
}
```

- Window is between the sequence numbers **a** and **c**
- **a** is considered earlier than **c**
- Checks if **b** is within the window

a=0	1	2	b=3	4	5	c=6
a=4	5	6	b=7	0	1	c=2
a=7	0	1	b=2	3	4	c=5

Go Back N

```
/* Protocol 5 (pipelining) allows multiple outstanding frames. The sender may transmit up to MAX_SEQ frames without waiting for an ack. In addition, unlike the previous protocols, the network layer is not assumed to have a new packet all the time. Instead, the network layer causes a network_layer_ready event when there is a packet to send. */
```

```
#define MAX_SEQ 7 /* should be  $2^n - 1$  */  
typedef enum {frame_arrival, cksum_err, timeout, network_layer_ready} event_type;  
#include "protocol.h"
```

```
static boolean between(seq_nr a, seq_nr b, seq_nr c)  
{  
/* Return true if  $a \leq b < c$  circularly; false otherwise. */  
if (((a <= b) && (b < c)) || ((c < a) && (a <= b)) || ((b < c) && (c < a)))  
    return(true);  
else  
    return(false);  
}
```

```
static void send_data(seq_nr frame_nr, seq_nr frame_expected, packet buffer[ ])  
{  
/* Construct and send a data frame. */  
    frame s; /* scratch variable */  
  
    s.info = buffer[frame_nr]; /* insert packet into frame */  
    s.seq = frame_nr; /* insert sequence number into frame */  
    s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1); /* piggyback ack */  
    to_physical_layer(&s); /* transmit the frame */  
    start_timer(frame_nr); /* start the timer running */  
}
```

- **Always** piggyback ACK with **every** data packet.
- This means that one side may continue to get ACK even though it is not sending any traffic.

- **s.ack** contains the sequence number of the last frame received
- Think of **s.ack** as **circular(frame_expected - 1)**
- Remember that $\text{circular}(x-1) = (x > 0) ? (x--) : \text{MAX_SEQ}$

- (Re-)Start a separate logical timer for every sent sequence number

Go Back N

```
void protocol5(void)
{
    seq_nr next_frame_to_send;          /* MAX_SEQ > 1; used for outbound stream */
    seq_nr ack_expected;                /* oldest frame as yet unacknowledged */
    seq_nr frame_expected;              /* next frame expected on inbound stream */
    frame r;                            /* scratch variable */
    packet buffer[MAX_SEQ + 1];         /* buffers for the outbound stream */
    seq_nr nbuffered;                   /* # output buffers currently in use */
    seq_nr i;                           /* used to index into the buffer array */
    event_type event;

    enable_network_layer();              /* allow network_layer_ready events */
    ack_expected = 0;                   /* next ack expected inbound */
    next_frame_to_send = 0;              /* next frame going out */
    frame_expected = 0;                 /* number of frame expected inbound */
    nbuffered = 0;                      /* initially no packets are buffered */
}
```

- Receiver window is of ***fixed size 1***
- Receiver window is between ***frame_expected*** and ***(frame_expected+1)***

Sender window is between ***ack_expected*** and ***next_frame_to_send***

Go Back N

```
while (true) {  
    wait_for_event(&event);          /* four possibilities: see event_type above */  
  
    switch(event) {  
        case network_layer_ready:    /* the network layer has a packet to send */  
            /* Accept, save, and transmit a new frame. */  
            from_network_layer(&buffer[next_frame_to_send]); /* fetch new packet */  
            nbuffered = nbuffered + 1; /* expand the sender's window */  
            send_data(next_frame_to_send, frame_expected, buffer); /* transmit the frame */  
            inc(next_frame_to_send); /* advance sender's upper window edge */  
            break;  
  
        case frame_arrival:          /* a data or control frame has arrived */  
            from_physical_layer(&r); /* get incoming frame from physical layer */  
  
            if (r.seq == frame_expected) {  
                /* Frames are accepted only in order. */  
                to_network_layer(&r.info); /* pass packet to network layer */  
                inc(frame_expected); /* advance lower edge of receiver's window */  
            }  
    }  
}
```

- We try to send packet in every loop
- We disable network layer when sender window is full
- Hence, we will always attempt to transmit the entire sender window

- Advance **upper end** of ***sender's*** window
⇒ increase sender's window

Go Back N

e.g. if Ack 3
then packets 2,1,0,...,
arrived.
called:
cumulative acknowledgement

```
/* Ack n implies n - 1, n - 2, etc. Check for this. */
while (between(ack_expected, r.ack, next_frame_to_send)) {
    /* Handle piggybacked ack. */
    nbuffered = nbuffered - 1; /* one frame fewer buffered */
    stop_timer(ack_expected); /* frame arrived intact; stop timer */
    inc(ack_expected); /* contract sender's window */
} break;

case cksum_err: break; /* just ignore bad frames */

case timeout: /* trouble; retransmit all outstanding frames */
    next_frame_to_send = ack_expected; /* start retransmitting here */
    for (i = 1; i <= nbuffered; i++) {
        send_data(next_frame_to_send, frame_expected, buffer); /* resend 1 frame */
        inc(next_frame_to_send); /* prepare to send the next one */
    }
}

if (nbuffered < MAX_SEQ)
    enable_network_layer();
else
    disable_network_layer();
}
```

- Outstanding packets are between **ack_expected** and **next_frame_to_send**

- Advance **lower end** of Sender's window until it hits *r.ack*
⇒ Reduce sender's window
⇒ Free more buffers

- Assume timeout because of frame loss
- Retransmit all frames **starting** from the **last ACKed** frame because receiver has window size one

- Max value of nbuffered is (MAX_SEQ)
- Maximum number of outstanding packets is MAX_SEQ **NOT** (MAX_SEQ+1)
 - Because the network layer is enabled only if nbuffered is strictly less than MAX_SEQ
- There are at most MAX_SEQ+1 distinct sequence numbers

Go Back N

- If there is **no reverse traffic**, protocol **fails**
 - sender will not receive ACK and not advance the window and it will keep timing-out and retransmitting → deadlock.
- Too many errors result in poor throughput because of retransmission.
- **No buffering** at the **receiver** since the receiver window size is 1.
- There is **buffering** at the **sender** since on timeout the sender re-transmits all outstanding frames.
- Need for *logically separate timer* per outstanding frame.
 - Expedite retransmission on loss (Optimization):
 - Sender sends packets 0→6 and puts timer on 6 (last set timer).
 - Suppose 0 is lost , 0 will be retransmitted only after timer of 6 expires.
 - Failure: (This can only happen if the protocol implementation is explicitly changed to allow: “setting the timer every n packets only”)
 - Suppose sender puts timer every 7 packets and that the NL has only 5 packets to send.
 - If any of the 5 packets gets lost, no retransmission since no timeout.

Go Back N

Sender window size (maximum outstanding frames) **must be** MAX_SEQ and **not** MAX_SEQ+1. I.e. if MAX_SEQ=7, we have 8 distinct sequence numbers (0→7) but sender window size=7(**Why??**)


- Assume sender window size=MAX_SEQ+1
- A transmits 0,1,2,3,4,5,6,7
- B received all of them and send ACK 7
- A sends another batch 0→7
- All 8 packets in the second batch are lost
- Now B has some data
- B sends data to A with ACK=7
- It is impossible for A to know whether the ACK is for the second batch or a duplicate ACK for the first batch

In the correct protocol case:

- A sends 0→6
- B sends ACK 6
- A sends a second batch 7,0,1,2,3,4,5
- Second batch is lost
- B sends data with ACK=6
- 6 is outside the sending window of A → A will ignore.

Go Back N: Efficient Timer Implementation

- Each frame times out independently of all the other ones. However, all of these timers can easily be simulated in software using a single hardware clock that causes interrupts periodically.
- The pending timeouts form a linked list, with each node of the list containing
 - the number of clock ticks until the timer expires.
 - the frame being timed.
 - a pointer to the next node.



(a) The queued timeouts.(b) The situation after the first timeout has expired.

- Initially, the real time is 10:00:00.000: three timeouts are pending, at 10:00:00.005, 10:00:00.013, and 10:00:00.019.
- Every time the hardware clock ticks, the real time is updated and the tick counter at the head of the list is decremented.
- When the tick counter becomes zero, a timeout is caused and the node is removed from the list.
- This organization requires the list to be scanned when *start timer* or *stop timer* is called, but it does not require much work per tick.